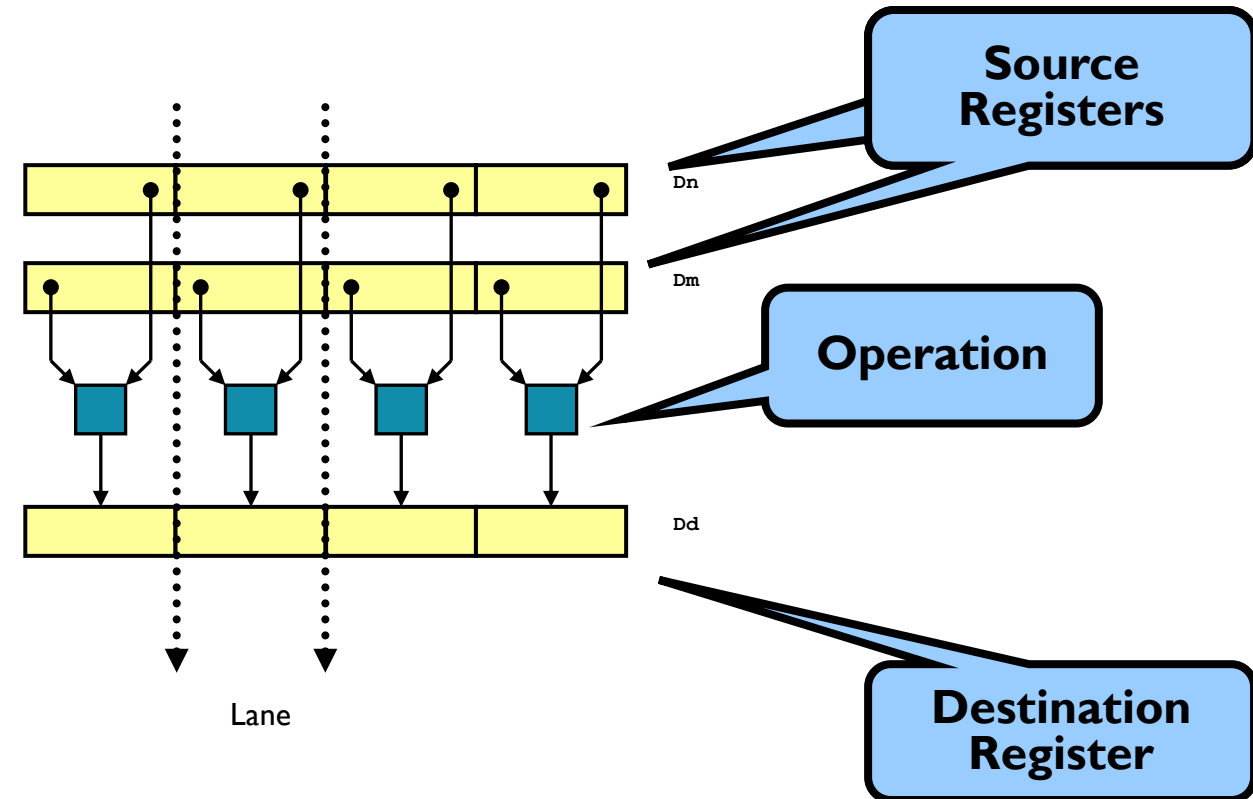


Moving your NEON optimizations to a 64-bit world

Ian Rickards
November 2014


What is NEON?

- NEON™ is a wide SIMD data processing architecture
 - Extension of the ARM® instruction set
 - 32 registers, 64-bit wide
(**AArch64: 32 registers**, 128-bit wide)
- NEON Instructions perform “Packed SIMD” processing
 - Registers are considered as vectors of elements of the same data type
 - Data types can be: signed/unsigned 8-bit, 16-bit, 32-bit, 64-bit, single precision float
(**AArch64: Double precision float**)
 - Instructions perform the same operation in all lanes




Changes in AArch64

- More registers
 - AArch32: 16x128-bit “Q-regs” Q0-Q15
 - AArch64: 32x128-bit “V-regs” V0-V31
- ‘Dual view’ no longer packed
 - 32 registers of each type: S, D, V
 - Clearer mapping of overlap
- Asm language changes
 - No ‘v’ in mnemonics
 - Width specifier moved to register description
 - 128-bit Q-regs renamed V-regs



vaddq.8 q0, q1, q2

8 = byte elements



add v0.8B, v1.8B, v2.8B

8 = Number of elements

B byte
H halfword (16b)
S single (32b)
D double (64b)

New capabilities in AArch64 'AdvancedSIMD'

- Reduction across all lanes
ADDV
- Rounding mode specified in instr
 - Not just 'round to nearest'
- Double precision
- Single destination register
VZIP, VSWAP now 2 instructions
- Crypto (ARMv8 but part of NEON)
- Ins instruction
- Table lookup larger table
- Saturating accumulate signed/unsigned
- AdvancedSIMD scalar
- REMOVED: high reg mapping, conditional execution

NEON use cases

- General purpose SIMD/DSP processing useful for many applications
- Support all new multimedia codecs



Watch any video in any format



Edit & Enhance captured videos
Video stabilization



Antialiased rendering & compositing



Advanced User Interfaces



Game processing



Process megapixel photos quickly



Voice recognition



Powerful multichannel hi-fi audio processing

NEON advantages

- **Easy to program**

- Clean vector architecture
- Off the shelf tools, OS support, commercial & opensource ecosystem support

- **Easy to debug**

- Single flow of control
- No separate DSP debugger

- **Fewer cycles needed**

- Neon will provide real-world 1.5x - 4x performance on typical video codecs
- Individual simple DSP algorithms can show larger performance boost (4x-8x)
- Provides overall **power saving** and **increased processing capabilities**
- **No overheads to 'calling' NEON**

How to use NEON

Automatically via OS

Opensource libraries

Vectorizing Compilers

- Uses NEON automatically from “C”

JIT compilers

- LLVM (e.g. Android Renderscript)

Commercial vendors

- e.g. commercial HEVC decoder

C Intrinsics

Assembler



Automatically via OS - NEON in Android


- Wide use of NEON optimizations in current Android source tree
- Many apps use NEON
 - Games (every game engine)
 - VR
 - Media editing / photo effects
 - Content creation



Component	ARMv7 NEON
VP8 (Google webm) decoder & encoder	YES (asm)
VP9 (Google webm) decoder & encoder	YES (asm)
JPEG	YES (asm)
Google WebP	YES (asm)
PNG decoder	YES (asm)
H.264 s/w decoder	YES (asm)
AMR WB encoder	YES (asm)
Skia	YES (asm)
WebRTC	YES (asm – FFT etc)
Renderscript	YES (via LLVM backend)
Blink (Chromium browser)	YES (intrinsics)

NEON optimizations in opensource

- **Google WebM** – **17,000** lines NEON code – both VP9 and VP8
- **Bluez** – official Linux Bluetooth protocol stack
- **Pixman** (part of cairo 2D graphics library)
- **ffmpeg (libav)** – **libavcodec** - LGPL media player
- **X264** - GPL H.264 encoder – can be used for video conferencing
- **Eigen2** – C++ vector math / linear algebra template library
- **Theorarm** – libtheora NEON version (optimized by Google)
- **Android libjpeg / libjpeg-turbo** – optimized JPEG decode
- **libpng** – NEON optimized PNG decode
- **FFTW** – NEON enabled FFT library
- **Liboil / liborc** – runtime compiler for SIMD processing
- **webkit** - used by Google Chrome browser
- **Ne10** – library of low-level optimized math routines
- **Cocos2d-x** – 2D game engine uses Ne10
- **Skia** – 2D graphics library
- **Android Renderscript**

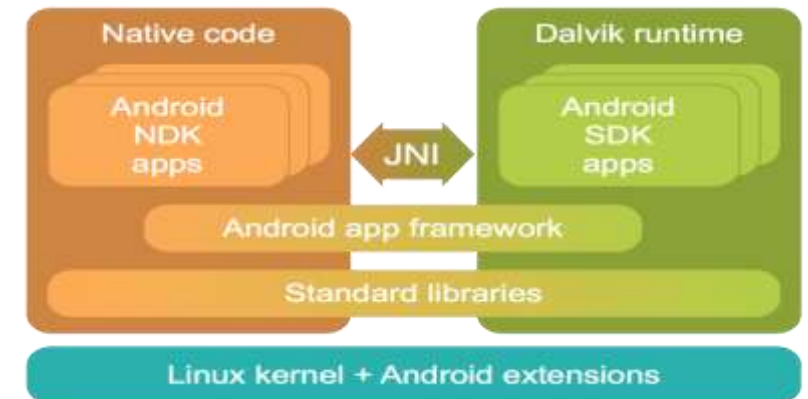


Android Native Development Kit (NDK) for ARM



- NDK is a toolkit to enable application developers to write native applications for the ARM processor
- NEON fully supported since NDK r5
- **64-bit support released in NDK r10 for “L”**

Android ABI	NEON support?
armeabi	No
armeabi-v7a	Optional - check cpu flags for NEON and ARMv8 crypto
arm64-v8a	Yes: NEON always present



Android™ applications can be written in Java, native ARM code, or a combination of the two

Using vectorizing compiler – gcc

```
int a[256], b[256], c[256];
```

```
foo () {  
    int i;
```

```
    for (i=0; i<256; i++){  
        a[i] = b[i] + c[i];  
    }  
}
```

AArch32

gcc -S -O3 -mcpu=cortex-a8
-mfpu=neon -mfloat-abi=softfp
test.c

.L2:

```
vldmia    r1!, {d18-d19}  
vldmia    r2!, {d16-d17}  
vadd.i32 q8, q9, q8  
vstmia    r3!, {d16-d17}  
cmp        r3, r0  
bne        .L2
```

AArch64

gcc -S -O3 test.c


.L2

```
ldr        q0, [x0], #16  
ldr        q1, [x2], #16  
cmp        x0, x3  
add       v0.4s, v0.4s, v1.4s  
str        q0, [x1], #16  
bne        .L2
```

gcc -ftree-vectorize is default at -O3
example built with linaro-gcc-4.9-2014.05 aarch64 release

What is vectorizing?


- Adding 16x 32-bit values (scalar)



- Using NEON using Q-register

AArch32


VADD
VADD
VADD
VADD




AArch64

ADD
ADD
ADD
ADD

VPADD



VPADD



Intrinsics

Fully compatible with AArch64

- **Include intrinsics header file (ACLE standard)**

```
#include <arm_neon.h>
```

- **Use special NEON data types which correspond to D and Q registers, e.g.**

<code>int8x8_t</code>	D-register 8x 8-bit values
<code>int16x4_t</code>	D-register 4x 16-bit values
<code>int32x4_t</code>	Q-register 4x 32-bit values

- **Use NEON intrinsics versions of instructions**

```
vin1 = vld1q_s32(ptr);  
vout = vaddq_s32(vin1, vin2);  
vst1q_s32(vout, ptr);
```

- **Strongly typed!**

- Use `vreinterpret_s16_s32()` to change the type

```
static inline void Filter_32_opaque_neon(unsigned x, unsigned y,  
                                          SkPMColor a00, SkPMColor a01,  
                                          SkPMColor a10, SkPMColor a11,  
                                          SkPMColor *dst) {  
    uint8x8_t vy, vconst16_8, v16_y, vres;  
    uint16x4_t vx, vconst16_16, v16_x, tmp;  
    uint32x2_t va0, va1;  
    uint16x8_t tmp1, tmp2;  
  
    vy = vdup_n_u8(y);                // duplicate y into vy  
    vconst16_8 = vmov_n_u8(16);        // set up constant in vcon  
    v16_y = vsub_u8(vconst16_8, vy);   // v16_y = 16-y  
  
    va0 = vdup_n_u32(a00);             // duplicate a00  
    va1 = vdup_n_u32(a10);             // duplicate a10  
    va0 = vset_lane_u32(a01, va0, 1);  // set top to a01  
    va1 = vset_lane_u32(a11, va1, 1);  // set top to a11  
  
    tmp1 = vmull_u8(vreinterpret_u8_u32(va0), v16_y); // tmp1 =  
    tmp2 = vmull_u8(vreinterpret_u8_u32(va1), vy);    // tmp2 =
```

NEON intrinsics

Pros

- Readability
- Reusability (inline functions, templates)
- Type checking (vreinterpret)
- Easier to debug
- **Portability to AArch64**
- Compiler can combine instructions (e.g. MAC)
- Compiler does register allocation
- Compiler does instruction scheduling

Cons

- Little control over registers used
- Does not always generate the code you expect


Compatibility

- C/intrinsics will port with no effort
- Asm requires reworking of .s file (mostly cosmetic, but can take advantage of additional registers)
- AArch64 NEON optimization in progress
 - ARM & Linaro working on key Android libraries using intrinsics
 - ffmpeg AArch64 NEON decoders (asm)
 - X264 AArch64 NEON encoder (asm)

AArch64 NEON coding technique	Compatible?
Vectorized “C”	Fully compatible
Intrinsics (“arm_neon.h”)	Fully compatible
Asm (.s)	Some porting required
Library routines	Yes, if library available

Fastest JPEG codecs: Android & libjpeg-turbo

- NEON optimizations integrated into
 - Official Android 4.4 (Kitkat) and later (plus partner-specific versions)
 - Libjpeg-turbo (opensource)
- Significantly improves speed of multi-megapixel image decode
- Benchmarked on 1.7GHz ARM Cortex®-A15
 - Optimized 34.4cycles/pix => 0.4s total for image
- Test image: 19.4Mpix
http://commons.wikimedia.org/wiki/File:Willaerts_Adam_The_Embarcation_of_the_Elector_Palantine_Oil_Canvas-huge.jpg




3268

5944

X264 – high quality H.264 encoding

- Can be used for highest-quality offline encoding – movies & tv content for on-demand services
- Full ARMv7 NEON optimizations
 - 5300 lines of NEON asm
- **New:** AArch64 NEON (Aug 2014)
- Performance results from Cortex-A15 processor @1.7GHz



File (media.xiph.org)	
park_joy_420_720p50.y4m	1280x720
sintel_trailer_2k_480p24.y4m	854x480

Wide range of NEON enabled low-cost dev boards

- Odroid XU3 \$179

- 4x 2.0GHz Cortex-A15 'Octa' b.L



- Cubieboard4 CC-A80

- 4x 2.0GHz Cortex-A15 + 4x 1.3GHz Cortex-A7



- Cubietruck \$65

- 4xCortex-A7



- Chromebook2

- 4xCortex-A15 'Octa' b.L



- **64-bit: ARM "Juno"**



- **64-bit: Nexus 9**



ARM

NEON summary

- NEON in AArch64 is much improved
 - More registers
 - New instructions
 - Cleaner instruction set
- Migrating to 64-bit
 - Use C or NEON intrinsics for best portability
 - Asm best in special circumstances, e.g. video codecs
Normally straightforward to port ARMv7 NEON to AArch64 NEON
 - NDK r10 provides full support – **start testing apps now!**
- Existing NEON documentation still very relevant
 - NEON Programmer's Guide
 - Blog entries
 - <http://www.arm.com/community/>
- Tune for AArch64
 - Extra registers
 - Double precision float
 - New instructions

Thank You

The trademarks featured in this presentation are registered and/or unregistered trademarks of ARM Limited (or its subsidiaries) in the EU and/or elsewhere. All rights reserved. Any other marks featured may be trademarks of their respective owners